

Lecture 8: Spatial data in R

Michael Lopez, Skidmore College

Data viz's in the news

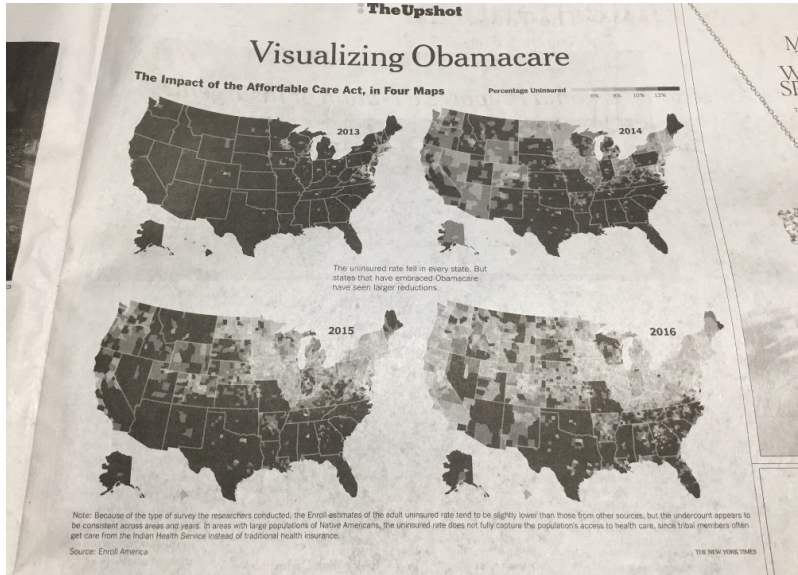


Figure 1: Uninsured over the last four years

Making Maps with R

Goals:

1. How to map in R
2. How to interpret maps in R
3. How to add to maps in R

Skills:

1. map package outlines
2. Conversion to data frames for ggplot2
3. ggmap to make maps
4. Mapping extras

Prerequisites

```
# packages we've used before
#install.packages(c("ggplot2", "dplyr", "stringr"))

# newer packages.
#install.packages(c("maps", "mapdata", "devtools"))

# the github version of ggmap
#devtools::install_github("dkahle/ggmap")
library(dplyr)
library(ggplot2)
library(ggmap)
library(maps)
library(mapdata)
```

- ▶ The maps package contains a lot of outlines of continents, countries, states, and counties that have been with R for a long time.
- ▶ Some examples: usa, nz, state, world, etc.
- ▶ The mapdata package contains a few more, higher-resolution outlines.
- ▶ The maps package comes with a plotting function, but, we will opt to use ggplot2 to plot the maps in the maps package.

Data frames from map outlines

- ▶ ggplot2 provides the `map_data()` function.
 - ▶ Think of it as a function that turns a series of points along an outline into a data frame of those points.
 - ▶ Syntax: `map_data("name")` where "name" is a quoted string of the name of a map in the `maps` or `mapdata` package
- ▶ USA map from `maps`:

```
usa <- map_data("usa")  
dim(usa)
```

```
## [1] 7243    6
```

```
usa %>% head(2)
```

```
##           long      lat group order region subregion  
## 1 -101.4078 29.74224     1     1   main      <NA>  
## 2 -101.3906 29.74224     1     2   main      <NA>
```

```
usa %>% tail(2)
```

```
##           long      lat group order      region subregion  
## 7251 -122.7104 48.21440    10   7251 whidbey island      <NA>  
## 7252 -122.6703 48.17429    10   7252 whidbey island      <NA>
```

Pacific Ocean

- ▶ Here is the high-res world map centered on the Pacific Ocean from mapdata

```
w2hr <- map_data("world2Hires")  
dim(w2hr)
```

```
## [1] 2274539      6
```

```
w2hr %>% head(3)
```

```
##      long      lat group order region subregion  
## 1 226.6336 58.42416     1     1  Canada      <NA>  
## 2 226.6314 58.42336     1     2  Canada      <NA>  
## 3 226.6122 58.41196     1     3  Canada      <NA>
```

```
w2hr %>% tail(3)
```


```
##      long      lat group  order      region subregion  
## 2276820 125.0100 11.15555 2284 2276820 Philippines  Leyte  
## 2276821 125.0111 11.14861 2284 2276821 Philippines  Leyte  
## 2276822 125.0155 11.13887 2284 2276822 Philippines  Leyte
```

Structure

- ▶ `long` is longitude. Points west of the prime meridian are negative.
- ▶ `lat` is latitude.
- ▶ `order` is how `ggplot` should “connect the dots”
- ▶ `region` and `subregion` is what region or subregion a set of points surrounds.
- ▶ `group` is whether adjacent points should be connected by lines
 - ▶ If they are in the same group, then they get connected
 - ▶ If in different groups, they don't
 - ▶ Equivalent to *lifting the pen* when connecting the dots

Plot the USA map

```
usa <- map_data("usa")
ggplot() +
  geom_polygon(data = usa, aes(x=long, y = lat, group = group)) +
  coord_fixed(1.3)
```




coord fixed

- ▶ Fixes relationship between one-unit in y direction and one unit in x direction
- ▶ Aspect ratio remains unchanged when saving plot
- ▶ Example: y unit 1.3 times longer than x unit
 - ▶ A different value might be needed closer to the poles.

Line and fill colors


- Fixed value of aesthetics go *outside* the aes function.

```
ggplot() +  
  geom_polygon(data = usa, aes(x=long, y = lat, group = group),  
              fill = NA, color = "red") +  
  coord_fixed(1.3)
```



Line and fill colors


```
ggplot() +  
  geom_polygon(data = usa, aes(x=long, y = lat, group = group),  
              fill = "violet", color = "blue") +  
  coord_fixed(1.3)
```



Adding points

- ▶ Let's add some points.


```
labs <- data.frame(  
  long = c(-73.7846, -74.0059),  
  lat = c(43.0831, 40.7128),  
  names = c("Toga", "NYC"))  
  
ggplot() +  
  geom_polygon(data = usa, aes(x=long, y = lat, group = group),  
    fill = "grey", color = "blue") +  
  geom_point(data = labs, aes(x = long, y = lat), color = "black", size = 2) +  
  coord_fixed(1.3)
```



The group aesthetic

Map without using the group aesthetic:

```
ggplot() +  
  geom_polygon(data = usa, aes(x=long, y = lat, group = group),  
              fill = "grey", color = "blue") +  
  geom_point(data = labs, aes(x = long, y = lat), color = "black", size = 2) +  
  coord_fixed(1.3)
```



State maps

We can also get a data frame of polygons that tell us about state boundaries:

```
states <- map_data("state")  
dim(states)
```

```
## [1] 15537      6
```


```
states %>% head(3)
```

```
##           long      lat group order  region subregion  
## 1 -87.46201 30.38968     1     1 alabama    <NA>  
## 2 -87.48493 30.37249     1     2 alabama    <NA>  
## 3 -87.52503 30.37249     1     3 alabama    <NA>
```

Plot all the states

- ▶ Set fill to region
- ▶ Lines of state borders are white.

```
ggplot(data = states) +  
  geom_polygon(aes(x = long, y = lat, fill = region, group = group),  
               color = "white") +  
  coord_fixed(1.3) +  
  guides(fill=FALSE) # do this to leave off the color legend
```




Subsetting states

```
west_coast <- states %>%  
  filter(region %in% c("california", "oregon", "washington"))  
west_coast %>% group_by(region) %>% count()
```

```
## # A tibble: 3 × 2  
##   region      n  
##   <chr> <int>  
## 1 california  516  
## 2    oregon   236  
## 3 washington 545
```


State graphs

```
ggplot(data = west_coast) +  
  geom_polygon(aes(x = long, y = lat, group = group),  
               fill = "palegreen", color = "black") +  
  coord_fixed(1.3)
```



County-level data

► New York data:

```
ny_df <- states %>% filter(region == "new york")
ny_df %>% head(3)
```

##	long	lat	group	order	region	subregion
## 1	-73.92874	40.80605	34	9050	new york	manhattan
## 2	-73.93448	40.78886	34	9051	new york	manhattan
## 3	-73.95166	40.77741	34	9052	new york	manhattan

► New York county lines

```
counties <- map_data("county")
ny_county <- counties %>% filter(region == "new york")
ny_county %>% head(3)
```

##	long	lat	group	order	region	subregion
## 1	-73.78550	42.46763	1795	52932	new york	albany
## 2	-74.25533	42.41034	1795	52933	new york	albany
## 3	-74.27252	42.41607	1795	52934	new york	albany

New York graph

- Note the use of `theme_map()` from `ggthemes` package


```
library(ggthemes)
ny_base <- ggplot(data = ny_df, mapping = aes(x = long, y = lat, group = group)) +
  coord_fixed(1.3) +
  geom_polygon(color = "black", fill = "gray")
ny_base + theme_map()
```



New York Graph with counties

- ▶ County boundaries in white
- ▶ ny_base

```
ny_base +  
  geom_polygon(data = ny_county, fill = NA, color = "white") +  
  geom_polygon(color = "black", fill = NA) +  
  theme_map()
```



County level data

- ▶ Hypothetical county level data for NY

```
unique.counties <- unique(ny_county$subregion)
n.counties <- length(unique.counties)
college.num <- 1:10
df.counties <- data.frame(subregion = unique.counties,
                          colleges = sample(1:10, n.counties, replace = TRUE))
df.counties %>% head(3)
```

```
##   subregion colleges
## 1    albany         6
## 2  allegany         3
## 3    bronx         1
```

Merging data together

- ▶ We now have the numbers that we want

```
ny_county %>% head(2)
```

```
##           long      lat group order  region subregion
## 1 -73.78550 42.46763  1795 52932 new york    albany
## 2 -74.25533 42.41034  1795 52933 new york    albany
```

```
df.counties %>% head(2)
```

```
##      subregion colleges
## 1      albany         6
## 2  allegany         3
```


- ▶ Merge together using `inner_join()`

```
ny_all <- inner_join(ny_county, df.counties, by = "subregion")
ny_all %>% sample_n(3)
```

```
##           long      lat group order  region  subregion colleges
## 1254 -75.44708 44.21515  1839 54229 new york st lawrence         2
## 1477 -73.41881 40.90919  1846 54459 new york    suffolk         10
## 805  -74.08917 42.95465  1823 53764 new york montgomery         8
```

NY college density


```
ny_base +  
  geom_polygon(data = ny_all, aes(fill = colleges), color = "white") +  
  geom_polygon(color = "black", fill = NA) +  
  coord_fixed(ratio = 1.3) +  
  theme_map()
```



NY college density, true zoom.

- Use the `xlim` and `ylim` arguments to `coord_fixed()`

```
ny_base +  
  geom_polygon(data = ny_all, aes(fill = colleges), color = "white") +  
  geom_polygon(color = "black", fill = NA) +  
  coord_fixed(xlim = c(-75, -73.0), ylim = c(42, 43), ratio = 1.3) +  
  theme_map()
```



How ggmap works

- ▶ Simplifies the process of downloading base maps from Google or Open Street Maps or Stamen Maps to use in the background of your plots.
- ▶ Sets axis scales
- ▶ Once you have gotten your maps, you make a call with `ggmap()` much as you would with `ggplot()`
- ▶ Let's do by example.

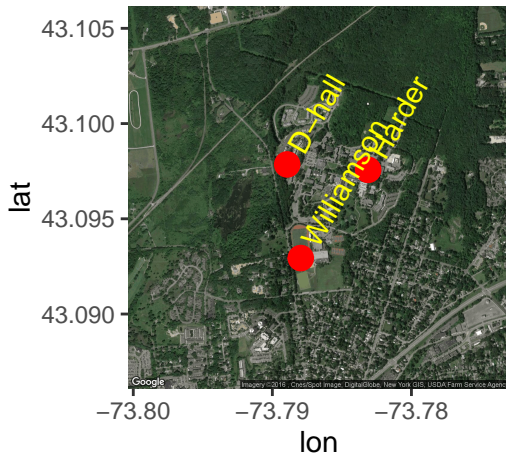
Skidmore data

- ▶ Here is a small data frame of points from Skidmore College
- ▶ note that ggmap tends to use “lon” instead of “long” for longitude

```
skid.df <- data.frame(name = c("Harder", "D-hall", "Williamson"),  
                      lon = c(-73.7831136, -73.7889339, -73.7879576),  
                      lat = c(43.0975717, 43.0978556, 43.0929376))
```

Google maps

```
# compute the mean lat and lon
ll_means <- colMeans(skid.df[2:3])
sq_map2 <- get_map(location = ll_means,
                    maptype = "satellite", source = "google", zoom = 15)
ggmap(sq_map2) +
  geom_point(data = skid.df, color = "red", size = 4) +
  geom_text(data = skid.df, aes(label = paste(" ", as.character(name), sep="")),
            angle = 60, hjust = 0, color = "yellow")
```



Google maps

```
sq_map3 <- get_map(location = ll_means,  
                   maptype = "terrain", source = "google", zoom = 15)  
ggmap(sq_map3) +  
  geom_point(data = skid.df, color = "red", size = 4) +  
  geom_text(data = skid.df, aes(label = paste(" ", as.character(name), sep="")),  
           angle = 60, hjust = 0, color = "black")
```

